

JWillis_D209_Data_Mining_PA2

January 8, 2023

0.1 D209 - Data Mining I - PA2

0.1.1 Import Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
from pandas import DataFrame
import sklearn.neighbors
from sklearn.neighbors import KNeighborsClassifier
from sklearn import preprocessing
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import train_test_split
```

0.1.2 Load Data From medical_clean.csv

```
[2]: # load data file
df = pd.read_csv('medical_clean.csv')
# quick test the data is present and see the shape
df.head()
```

```
[2]: CaseOrder Customer_id Interaction \
0      1      C412403  8cd49b13-f45a-4b47-a2bd-173ffa932c2f
1      2      Z919181  d2450b70-0337-4406-bdbb-bc1037f1734c
2      3      F995323  a2057123-abf5-4a2c-abad-8ffe33512562
3      4      A879973  1dec528d-eb34-4079-adce-0d7a40e82205
4      5      C544523  5885f56b-d6da-43a3-8760-83583af94266

      UID      City State      County      Zip \
0  3a83ddb66e2ae73798bdf1d705dc0932      Eva      AL      Morgan  35621
1  176354c5eef714957d486009feabf195      Marianna      FL      Jackson  32446
2  e19a0fa00aeda885b8a436757e889bc9      Sioux Falls      SD      Minnehaha  57110
3  cd17d7b6d152cb6f23957346d11c3f07      New Richland      MN      Waseca  56072
4  d2f0425877b10ed6bb381f3e2579424a      West Point      VA      King William  23181

      Lat      Lng ... TotalCharge Additional_charges Item1 Item2 Item3 \
```

0	34.34960	-86.72508	...	3726.702860	17939.403420	3	3	2
1	30.84513	-85.22907	...	4193.190458	17612.998120	3	4	3
2	43.54321	-96.63772	...	2434.234222	17505.192460	2	4	4
3	43.89744	-93.51479	...	2127.830423	12993.437350	3	5	5
4	37.59894	-76.88958	...	2113.073274	3716.525786	2	1	3

	Item4	Item5	Item6	Item7	Item8
0	2	4	3	3	4
1	4	4	4	3	3
2	4	3	4	3	3
3	3	4	5	5	5
4	3	5	3	4	3

[5 rows x 50 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CaseOrder              10000 non-null  int64
1   Customer_id            10000 non-null  object
2   Interaction             10000 non-null  object
3   UID                    10000 non-null  object
4   City                   10000 non-null  object
5   State                  10000 non-null  object
6   County                 10000 non-null  object
7   Zip                    10000 non-null  int64
8   Lat                    10000 non-null  float64
9   Lng                    10000 non-null  float64
10  Population              10000 non-null  int64
11  Area                    10000 non-null  object
12  TimeZone                10000 non-null  object
13  Job                     10000 non-null  object
14  Children                10000 non-null  int64
15  Age                     10000 non-null  int64
16  Income                  10000 non-null  float64
17  Marital                 10000 non-null  object
18  Gender                  10000 non-null  object
19  ReAdmis                 10000 non-null  object
20  VitD_levels             10000 non-null  float64
21  Doc_visits              10000 non-null  int64
22  Full_meals_eaten        10000 non-null  int64
23  vitD_supp               10000 non-null  int64
24  Soft_drink              10000 non-null  object
```

```

25 Initial_admin      10000 non-null object
26 HighBlood          10000 non-null object
27 Stroke              10000 non-null object
28 Complication_risk  10000 non-null object
29 Overweight          10000 non-null object
30 Arthritis           10000 non-null object
31 Diabetes            10000 non-null object
32 Hyperlipidemia      10000 non-null object
33 BackPain            10000 non-null object
34 Anxiety             10000 non-null object
35 Allergic_rhinitis   10000 non-null object
36 Reflux_esophagitis  10000 non-null object
37 Asthma              10000 non-null object
38 Services            10000 non-null object
39 Initial_days        10000 non-null float64
40 TotalCharge         10000 non-null float64
41 Additional_charges  10000 non-null float64
42 Item1               10000 non-null int64
43 Item2               10000 non-null int64
44 Item3               10000 non-null int64
45 Item4               10000 non-null int64
46 Item5               10000 non-null int64
47 Item6               10000 non-null int64
48 Item7               10000 non-null int64
49 Item8               10000 non-null int64
dtypes: float64(7), int64(16), object(27)
memory usage: 3.8+ MB


```

Look for Missing Values

```

[4]: # Mapping to view missing data...none present.
sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap='viridis');

```



```
[5]: df.describe()
```

```
[5]:
```

	CaseOrder	Zip	Lat	Lng	Population \
count	10000.00000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.50000	50159.323900	38.751099	-91.243080	9965.253800
std	2886.89568	27469.588208	5.403085	15.205998	14824.758614
min	1.00000	610.000000	17.967190	-174.209700	0.000000
25%	2500.75000	27592.000000	35.255120	-97.352982	694.750000
50%	5000.50000	50207.000000	39.419355	-88.397230	2769.000000
75%	7500.25000	72411.750000	42.044175	-80.438050	13945.000000
max	10000.00000	99929.000000	70.560990	-65.290170	122814.000000

	Children	Age	Income	VitD_levels	Doc_visits \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	2.097200	53.511700	40490.495160	17.964262	5.012200
std	2.163659	20.638538	28521.153293	2.017231	1.045734

min	0.000000	18.000000	154.080000	9.806483	1.000000
25%	0.000000	36.000000	19598.775000	16.626439	4.000000
50%	1.000000	53.000000	33768.420000	17.951122	5.000000
75%	3.000000	71.000000	54296.402500	19.347963	6.000000
max	10.000000	89.000000	207249.100000	26.394449	9.000000

	...	TotalCharge	Additional_charges	Item1	Item2 \
count	...	10000.000000	10000.000000	10000.000000	10000.000000
mean	...	5312.172769	12934.528587	3.518800	3.506700
std	...	2180.393838	6542.601544	1.031966	1.034825
min	...	1938.312067	3125.703000	1.000000	1.000000
25%	...	3179.374015	7986.487755	3.000000	3.000000
50%	...	5213.952000	11573.977735	4.000000	3.000000
75%	...	7459.699750	15626.490000	4.000000	4.000000
max	...	9180.728000	30566.070000	8.000000	7.000000

		Item3	Item4	Item5	Item6	Item7 \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean		3.511100	3.515100	3.496900	3.522500	3.494000
std		1.032755	1.036282	1.030192	1.032376	1.021405
min		1.000000	1.000000	1.000000	1.000000	1.000000
25%		3.000000	3.000000	3.000000	3.000000	3.000000
50%		4.000000	4.000000	3.000000	4.000000	3.000000
75%		4.000000	4.000000	4.000000	4.000000	4.000000
max		8.000000	7.000000	7.000000	7.000000	7.000000

	Item8
count	10000.000000
mean	3.509700
std	1.042312
min	1.000000
25%	3.000000
50%	3.000000
75%	4.000000
max	7.000000

[8 rows x 23 columns]

0.1.3 Describe and Explore Numeric Fields:

```
[6]: df.describe(include = [np.number])
```

[6]:	CaseOrder	Zip	Lat	Lng	Population \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.500000	50159.323900	38.751099	-91.243080	9965.253800
std	2886.89568	27469.588208	5.403085	15.205998	14824.758614
min	1.000000	610.000000	17.967190	-174.209700	0.000000

25%	2500.75000	27592.000000	35.255120	-97.352982	694.750000
50%	5000.50000	50207.000000	39.419355	-88.397230	2769.000000
75%	7500.25000	72411.750000	42.044175	-80.438050	13945.000000
max	10000.00000	99929.000000	70.560990	-65.290170	122814.000000

	Children	Age	Income	VitD_levels	Doc_visits \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	2.097200	53.511700	40490.495160	17.964262	5.012200
std	2.163659	20.638538	28521.153293	2.017231	1.045734
min	0.000000	18.000000	154.080000	9.806483	1.000000
25%	0.000000	36.000000	19598.775000	16.626439	4.000000
50%	1.000000	53.000000	33768.420000	17.951122	5.000000
75%	3.000000	71.000000	54296.402500	19.347963	6.000000
max	10.000000	89.000000	207249.100000	26.394449	9.000000

	...	TotalCharge	Additional_charges	Item1	Item2 \
count	...	10000.000000	10000.000000	10000.000000	10000.000000
mean	...	5312.172769	12934.528587	3.518800	3.506700
std	...	2180.393838	6542.601544	1.031966	1.034825
min	...	1938.312067	3125.703000	1.000000	1.000000
25%	...	3179.374015	7986.487755	3.000000	3.000000
50%	...	5213.952000	11573.977735	4.000000	3.000000
75%	...	7459.699750	15626.490000	4.000000	4.000000
max	...	9180.728000	30566.070000	8.000000	7.000000

	Item3	Item4	Item5	Item6	Item7 \
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	3.511100	3.515100	3.496900	3.522500	3.494000
std	1.032755	1.036282	1.030192	1.032376	1.021405
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000	3.000000	3.000000
50%	4.000000	4.000000	3.000000	4.000000	3.000000
75%	4.000000	4.000000	4.000000	4.000000	4.000000
max	8.000000	7.000000	7.000000	7.000000	7.000000

	Item8
count	10000.000000
mean	3.509700
std	1.042312
min	1.000000
25%	3.000000
50%	3.000000
75%	4.000000
max	7.000000

[8 rows x 23 columns]

```
[7]: df_num = df.select_dtypes(include='number')
df_num.head()
```

```
[7]: CaseOrder    Zip      Lat      Lng  Population  Children  Age  Income \
0           1  35621  34.34960 -86.72508         2951         1   53  86575.93
1           2  32446  30.84513 -85.22907         11303         3   51  46805.99
2           3  57110  43.54321 -96.63772         17125         3   53  14370.14
3           4  56072  43.89744 -93.51479          2162         0   78  39741.49
4           5  23181  37.59894 -76.88958          5287         1   22   1209.56
```

```
      VitD_levels  Doc_visits  ...  TotalCharge  Additional_charges  Item1 \
0      19.141466           6  ...  3726.702860         17939.403420         3
1      18.940352           4  ...  4193.190458         17612.998120         3
2      18.057507           4  ...  2434.234222         17505.192460         2
3      16.576858           4  ...  2127.830423         12993.437350         3
4      17.439069           5  ...  2113.073274          3716.525786         2
```

```
      Item2  Item3  Item4  Item5  Item6  Item7  Item8
0         3     2     2     4     3     3     4
1         4     3     4     4     4     3     3
2         4     4     4     3     4     3     3
3         5     5     3     4     5     5     5
4         1     3     3     5     3     4     3
```

[5 rows x 23 columns]

0.1.4 Describe and Explore Categorical Fields:

```
[8]: df.describe(exclude = [np.number])
```

```
[8]: Customer_id      Interaction \
count      10000      10000
unique      10000      10000
top      C412403  8cd49b13-f45a-4b47-a2bd-173ffa932c2f
freq           1           1

      UID      City  State      County      Area \
count      10000  10000  10000      10000  10000
unique      10000   6072    52       1607     3
top      3a83ddb66e2ae73798bdf1d705dc0932  Houston    TX  Jefferson  Rural
freq           1     36   553       118   3369

      TimeZone      Job  Marital  ... \
count      10000      10000  10000  ...
unique         26       639     5  ...
top  America/New_York  Outdoor activities/education manager  Widowed  ...
freq      3889           29    2045  ...
```

	Overweight	Arthritis	Diabetes	Hyperlipidemia	BackPain	Anxiety	\
count	10000	10000	10000	10000	10000	10000	
unique	2	2	2	2	2	2	
top	Yes	No	No	No	No	No	
freq	7094	6426	7262	6628	5886	6785	

	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services
count	10000	10000	10000	10000
unique	2	2	2	4
top	No	No	No	Blood Work
freq	6059	5865	7107	5265

[4 rows x 27 columns]

```
[9]: df_cat = df.select_dtypes(exclude='number')
df_cat.head()
```

```
[9]: Customer_id      Interaction \
0      C412403  8cd49b13-f45a-4b47-a2bd-173ffa932c2f
1      Z919181  d2450b70-0337-4406-bdbb-bc1037f1734c
2      F995323  a2057123-abf5-4a2c-abad-8ffe33512562
3      A879973  1dec528d-eb34-4079-adce-0d7a40e82205
4      C544523  5885f56b-d6da-43a3-8760-83583af94266
```

	UID	City	State	County	\
0	3a83ddb66e2ae73798bdf1d705dc0932	Eva	AL	Morgan	
1	176354c5eef714957d486009feabf195	Marianna	FL	Jackson	
2	e19a0fa00aeda885b8a436757e889bc9	Sioux Falls	SD	Minnehaha	
3	cd17d7b6d152cb6f23957346d11c3f07	New Richland	MN	Waseca	
4	d2f0425877b10ed6bb381f3e2579424a	West Point	VA	King William	

	Area	TimeZone	Job	Marital	\
0	Suburban	America/Chicago	Psychologist, sport and exercise	Divorced	
1	Urban	America/Chicago	Community development worker	Married	
2	Suburban	America/Chicago	Chief Executive Officer	Widowed	
3	Suburban	America/Chicago	Early years teacher	Married	
4	Rural	America/New_York	Health promotion specialist	Widowed	

	...	Overweight	Arthritis	Diabetes	Hyperlipidemia	BackPain	Anxiety	\
0	...	No	Yes	Yes	No	Yes	Yes	
1	...	Yes	No	No	No	No	No	
2	...	Yes	No	Yes	No	No	No	
3	...	No	Yes	No	No	No	No	
4	...	No	No	No	Yes	No	No	

	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services
--	-------------------	--------------------	--------	----------

0	Yes	No	Yes	Blood Work
1	No	Yes	No	Intravenous
2	No	No	No	Blood Work
3	No	Yes	Yes	Blood Work
4	Yes	No	No	CT Scan

[5 rows x 27 columns]

```
[10]: df[['ReAdmis']].describe()
df
```

```
[10]: CaseOrder Customer_id Interaction \
0      1      C412403  8cd49b13-f45a-4b47-a2bd-173ffa932c2f
1      2      Z919181  d2450b70-0337-4406-bdbb-bc1037f1734c
2      3      F995323  a2057123-abf5-4a2c-abad-8ffe33512562
3      4      A879973  1dec528d-eb34-4079-adce-0d7a40e82205
4      5      C544523  5885f56b-d6da-43a3-8760-83583af94266
...
9995   9996   B863060  a25b594d-0328-486f-a9b9-0567eb0f9723
9996   9997   P712040  70711574-f7b1-4a17-b15f-48c54564b70f
9997   9998   R778890  1d79569d-8e0f-4180-a207-d67ee4527d26
9998   9999   E344109  f5a68e69-2a60-409b-a92f-ac0847b27db0
9999  10000   I569847  bc482c02-f8c9-4423-99de-3db5e62a18d5

      UID      City State      County \
0  3a83ddb66e2ae73798bdf1d705dc0932      Eva      AL      Morgan
1  176354c5eef714957d486009feabf195      Marianna      FL      Jackson
2  e19a0fa00aeda885b8a436757e889bc9      Sioux Falls      SD      Minnehaha
3  cd17d7b6d152cb6f23957346d11c3f07      New Richland      MN      Waseca
4  d2f0425877b10ed6bb381f3e2579424a      West Point      VA      King William
...
9995  39184dc28cc038871912ccc450049e5      Norlina      NC      Warren
9996  3cd124ccd43147404292e883bf9ec55c      Milmay      NJ      Atlantic
9997  41b770aeee97a5b9e7f69c906a8119d7      Southside      TN      Montgomery
9998  2bb491ef5b1beb1fed758cc6885c167a      Quinn      SD      Pennington
9999  95663a202338000abdf7e09311c2a8a1      Coraopolis      PA      Allegheny

      Zip      Lat      Lng ... TotalCharge Additional_charges Item1 \
0  35621  34.34960 -86.72508 ... 3726.702860      17939.403420      3
1  32446  30.84513 -85.22907 ... 4193.190458      17612.998120      3
2  57110  43.54321 -96.63772 ... 2434.234222      17505.192460      2
3  56072  43.89744 -93.51479 ... 2127.830423      12993.437350      3
4  23181  37.59894 -76.88958 ... 2113.073274      3716.525786      2
...
9995  27563  36.42886 -78.23716 ... 6850.942000      8927.642000      3
9996   8340  39.43609 -74.87302 ... 7741.690000      28507.150000      3
9997  37171  36.36655 -87.29988 ... 8276.481000      15281.210000      3
```

9998	57775	44.10354	-102.01590	...	7644.483000	7781.678000	5
9999	15108	40.49998	-80.19959	...	7887.553000	11643.190000	4

	Item2	Item3	Item4	Item5	Item6	Item7	Item8
0	3	2	2	4	3	3	4
1	4	3	4	4	4	3	3
2	4	4	4	3	4	3	3
3	5	5	3	4	5	5	5
4	1	3	3	5	3	4	3
...
9995	2	2	3	4	3	4	2
9996	3	4	2	5	3	4	4
9997	3	3	4	4	2	3	2
9998	5	3	4	4	3	4	3
9999	3	3	2	3	6	4	3

[10000 rows x 50 columns]

```
[11]: df[['Gender', 'ReAdmis']].head()
```

```
[11]:   Gender ReAdmis
0    Male      No
1  Female      No
2  Female      No
3    Male      No
4  Female      No
```

```
[12]: df_2num = df.copy()
      df_2num.columns
```

```
[12]: Index(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City', 'State',
            'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area', 'TimeZone', 'Job',
            'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis',
            'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp',
            'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke',
            'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes',
            'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis',
            'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',
            'TotalCharge', 'Additional_charges', 'Item1', 'Item2', 'Item3', 'Item4',
            'Item5', 'Item6', 'Item7', 'Item8'],
            dtype='object')
```

0.1.5 Transform Catagorical Data To Numerical

```
[13]: def cat_to_num(series):  
        series = series.astype('category')  
        return series.cat.codes  
  
df_2num = df_2num.apply(cat_to_num) # Male = 1, ReAdmis Yes = 1  
df_2num.head()
```

```
[13]:
```

	CaseOrder	Customer_id	Interaction	UID	City	State	County	Zip	Lat	\
0	0	887	5453	2240	1675	1	975	2952	1774	
1	1	9960	8124	862	3232	9	712	2683	662	
2	2	2368	6269	8801	5021	42	951	4944	7305	
3	3	320	1167	7953	3782	23	1514	4809	7443	
4	4	943	3430	8170	5814	46	766	1778	3174	

	Lng	...	TotalCharge	Additional_charges	Item1	Item2	Item3	Item4	\
0	4851	...	3949	7427	2	2	1	1	
1	5178	...	4620	7363	2	3	2	3	
2	2470	...	384	7337	1	3	3	3	
3	3245	...	37	5566	2	4	4	2	
4	7357	...	28	104	1	0	2	2	

	Item5	Item6	Item7	Item8
0	3	2	2	3
1	3	3	2	2
2	2	3	2	2
3	3	4	4	4
4	4	2	3	2

[5 rows x 50 columns]

```
[14]: df_2num.shape
```

```
[14]: (10000, 50)
```

```
[15]: df[['Gender', 'ReAdmis']].head()
```

```
[15]:
```

	Gender	ReAdmis
0	Male	No
1	Female	No
2	Female	No
3	Male	No
4	Female	No

```
[16]: #df['ReAdmis_Yes'] = df['ReAdmis']  
      #df['ReAdmis_Yes'] = df['ReAdmis'].eq('Yes').astype(int)
```

```
#df['ReAdmis_Yes']
```

```
[17]: df_2num.columns
```

```
[17]: Index(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City', 'State',  
        'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area', 'TimeZone', 'Job',  
        'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis',  
        'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp',  
        'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke',  
        'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes',  
        'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis',  
        'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',  
        'TotalCharge', 'Additional_charges', 'Item1', 'Item2', 'Item3', 'Item4',  
        'Item5', 'Item6', 'Item7', 'Item8'],  
        dtype='object')
```

```
[18]: df['Gender'].value_counts()
```

```
[18]: Female      5018  
      Male      4768  
      Nonbinary   214  
      Name: Gender, dtype: int64
```

```
[19]: df_2num['Gender'].value_counts() # why is female zero? Alphabetical sequence?
```

```
[19]: 0      5018  
      1      4768  
      2       214  
      Name: Gender, dtype: int64
```

```
[20]: df['Initial_admin'].value_counts()
```

```
[20]: Emergency Admission      5060  
      Elective Admission    2504  
      Observation Admission  2436  
      Name: Initial_admin, dtype: int64
```

```
[21]: df_2num['Initial_admin'].value_counts() # Seems to be ordered alphabetically.
```

```
[21]: 1      5060  
      0      2504  
      2      2436  
      Name: Initial_admin, dtype: int64
```

```
[22]: df['Initial_admin'].head()
```

```
[22]: 0    Emergency Admission
      1    Emergency Admission
      2    Elective Admission
      3    Elective Admission
      4    Elective Admission
      Name: Initial_admin, dtype: object
```

0.1.6 Data Preparation

0.1.7 Prune Data

```
[23]: df_2num.columns
```

```
[23]: Index(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City', 'State',
          'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area', 'TimeZone', 'Job',
          'Children', 'Age', 'Income', 'Marital', 'Gender', 'ReAdmis',
          'VitD_levels', 'Doc_visits', 'Full_meals_eaten', 'vitD_supp',
          'Soft_drink', 'Initial_admin', 'HighBlood', 'Stroke',
          'Complication_risk', 'Overweight', 'Arthritis', 'Diabetes',
          'Hyperlipidemia', 'BackPain', 'Anxiety', 'Allergic_rhinitis',
          'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',
          'TotalCharge', 'Additional_charges', 'Item1', 'Item2', 'Item3', 'Item4',
          'Item5', 'Item6', 'Item7', 'Item8'],
          dtype='object')
```

```
[24]: df_2num['Interaction'].head()
```

```
[24]: 0    5453
      1    8124
      2    6269
      3    1167
      4    3430
      Name: Interaction, dtype: int16
```

```
[25]: # Start pruning non-relevant series
pruned_df = df_2num.drop(['CaseOrder', 'Customer_id', 'Interaction', 'UID',
↪ 'City', 'State', \
                        'County', 'Zip', 'Lat', 'Lng', 'Population', 'Area',
↪ 'TimeZone', 'Job', \
                        'Children', 'Income', 'Marital', 'Full_meals_eaten',
↪ 'Soft_drink', \
                        'Item1', 'Item2', 'Item3', 'Item4', 'Item5', 'Item6',
↪ 'Item7', 'Item8'], axis=1)
pruned_df.head()
```

```
[25]:   Age  Gender  ReAdmis  VitD_levels  Doc_visits  vitD_supp  Initial_admin \
0    35        1         0         7175          5           0           1
```

1	33	0	0	6885	3	1	1
2	35	0	0	5219	3	0	0
3	60	1	0	2422	3	0	0
4	4	0	0	3978	4	2	0

	HighBlood	Stroke	Complication_risk	...	Hyperlipidemia	BackPain	\
0	1	0	2	...	0	1	
1	1	0	0	...	0	0	
2	1	0	2	...	0	0	
3	0	1	2	...	0	0	
4	0	0	1	...	1	0	

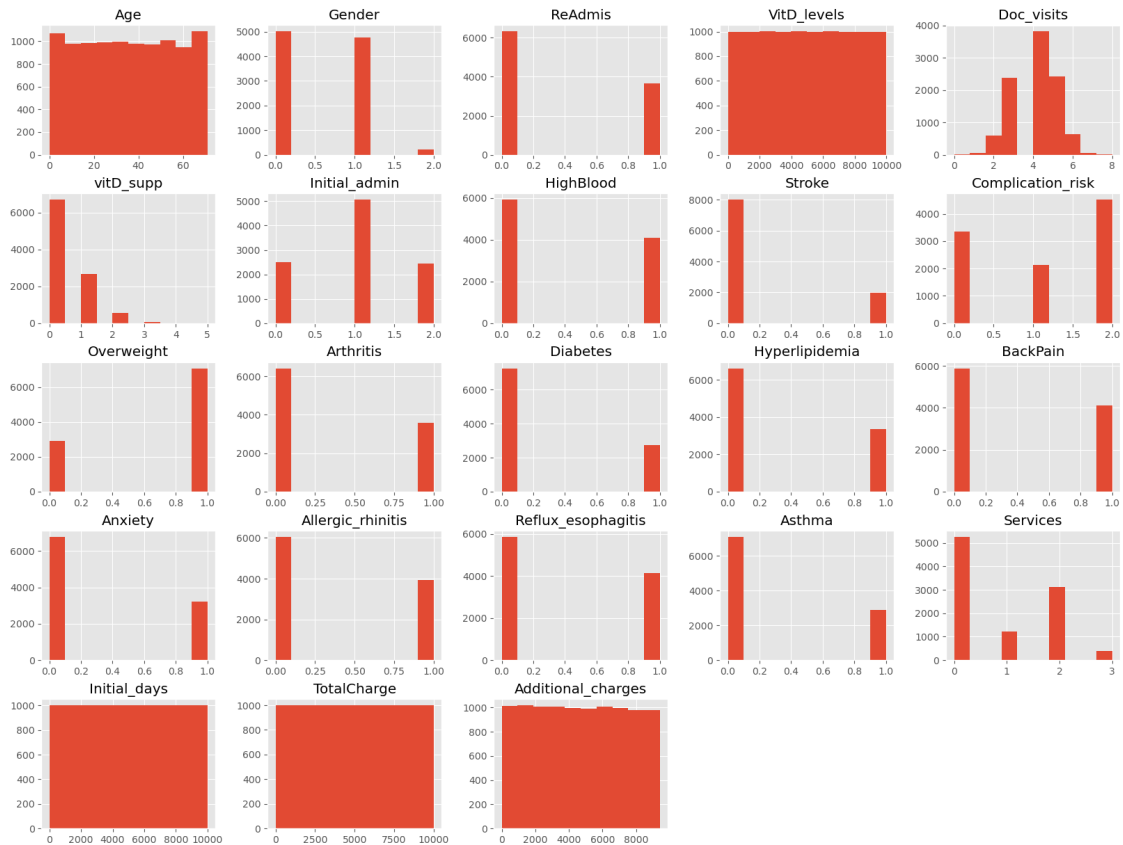
	Anxiety	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services	\
0	1	1	0	1	0	
1	0	0	1	0	2	
2	0	0	0	0	0	
3	0	0	1	1	0	
4	0	1	0	0	1	

	Initial_days	TotalCharge	Additional_charges
0	3234	3949	7427
1	4128	4620	7363
2	1397	384	7337
3	285	37	5566
4	101	28	104

[5 rows x 23 columns]

```
[26]: # https://www.datacamp.com/community/tutorials/
      ↪ preprocessing-in-data-science-part-1-centering-scaling-and-knn
plt.style.use('ggplot')
# df = pd.read_csv('http://archive.ics.uci.edu/ml/machine-learning-databases/
      ↪ wine-quality/winequality-red.csv ', sep = ';')
X = pruned_df.drop('ReAdmis', 1).values # drop target variable
y1 = pruned_df['ReAdmis'].values
pd.DataFrame.hist(pruned_df, figsize = [20,15]);
```

```
/var/folders/45/_087y05165x0c7wb_dw4k6nh0000gn/T/ipykernel_62417/922365799.py:4:
FutureWarning: In a future version of pandas all arguments of DataFrame.drop
except for the argument 'labels' will be keyword-only.
X = pruned_df.drop('ReAdmis', 1).values # drop target variable
```



```
[27]: # https://realpython.com/knn-python/
# Correlations with target?
#correlation_matrix = pruned_df.corr()

#print(correlation_matrix["ReAdmis_Yes"] > 0.5)
```

0.1.8 Split the Training Data

```
[28]: # Split dataset into two data groups: train and test
from sklearn.model_selection import train_test_split
train, test = train_test_split(pruned_df, test_size = 0.2) # 80% of data
↳ assigned to train
```

```
[29]: print("Train Shape: ", train.shape)
print("Test Shape: ", test.shape)
```

Train Shape: (8000, 23)

Test Shape: (2000, 23)

```
[30]: # Notice the first column (row numbers) shows the choices are quite random
print("Train: ", train.head())
print('-----'*10)
print("Test: ", test.head())
```

Train:	Age	Gender	ReAdmis	VitD_levels	Doc_visits	vitD_supp
Initial_admin \						
7141	29	0	1	9734	4	2
7618	15	1	1	119	2	0
6254	49	0	1	7922	4	0
7265	16	1	1	1539	4	0
2514	7	0	0	8964	4	0

	HighBlood	Stroke	Complication_risk	...	Hyperlipidemia	BackPain	\
7141	1	1	2	...	0	0	
7618	0	0	2	...	0	1	
6254	0	0	0	...	1	0	
7265	0	0	2	...	0	0	
2514	1	0	2	...	0	1	

	Anxiety	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services	\
7141	0	0	0	0	2	
7618	0	0	1	0	0	
6254	1	0	1	0	0	
7265	0	0	1	0	0	
2514	1	1	0	0	0	

	Initial_days	TotalCharge	Additional_charges
7141	9556	9290	7162
7618	8985	9000	1167
6254	8729	8953	4634
7265	7038	6051	1040
2514	2795	3494	2893

[5 rows x 23 columns]

```
-----
```

Test:	Age	Gender	ReAdmis	VitD_levels	Doc_visits	vitD_supp
Initial_admin \						
4184	68	0	0	3245	5	0
9309	33	1	0	9710	3	0
4963	6	1	0	9499	6	1
2320	26	0	0	154	4	0
8161	65	0	1	4176	4	0

	HighBlood	Stroke	Complication_risk	...	Hyperlipidemia	BackPain	\
4184	1	1	0	...	0	1	

9309	1	0	2	...	0	1
4963	0	0	0	...	1	0
2320	1	0	0	...	1	1
8161	0	0	1	...	0	0

	Anxiety	Allergic_rhinitis	Reflux_esophagitis	Asthma	Services	\
4184	0	0	0	0	0	
9309	0	1	1	0	0	
4963	1	1	1	0	3	
2320	0	0	1	0	2	
8161	0	0	0	0	1	

	Initial_days	TotalCharge	Additional_charges
4184	4893	4965	9379
9309	6136	5870	7254
4963	4837	4962	713
2320	4591	4591	6775
8161	8203	8141	6115

[5 rows x 23 columns]

0.1.9 Build Decision Tree

```
[31]: pruned_df.columns
```

```
[31]: Index(['Age', 'Gender', 'ReAdmis', 'VitD_levels', 'Doc_visits', 'vitD_supp',
            'Initial_admin', 'HighBlood', 'Stroke', 'Complication_risk',
            'Overweight', 'Arthritis', 'Diabetes', 'Hyperlipidemia', 'BackPain',
            'Anxiety', 'Allergic_rhinitis', 'Reflux_esophagitis', 'Asthma',
            'Services', 'Initial_days', 'TotalCharge', 'Additional_charges'],
            dtype='object')
```

```
[32]: # Feed training data into ML Algorithm and Build Decision Tree
from sklearn.tree import DecisionTreeClassifier
clf=DecisionTreeClassifier(max_leaf_nodes=10) # Tried default, 5, 10, 15; 10 &
↪15 were the same
clf=clf.fit(train[['Age', 'Gender', 'VitD_levels', 'Doc_visits', 'vitD_supp',
↪'Initial_admin', \
                                'HighBlood', 'Stroke', 'Complication_risk', 'Overweight',
↪'Arthritis', \
                                'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety',
↪'Allergic_rhinitis', \
                                'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',
↪'TotalCharge', \
                                'Additional_charges']], train['ReAdmis'])
```

```
[33]: clf.feature_importances_
```

```
[33]: array([0.          , 0.          , 0.00188001, 0.          , 0.          ,
         0.          , 0.          , 0.00102016, 0.          , 0.          ,
         0.          , 0.          , 0.          , 0.          , 0.          ,
         0.          , 0.          , 0.00117126, 0.          , 0.98990247,
         0.00602611, 0.          ])
```

```
[34]: # https://stackoverflow.com/questions/1494492/
      ↪ graphviz-how-to-go-from-dot-to-a-graph
from sklearn import tree
with open("med_data.gif", 'w') as f:
    f = tree.export_graphviz(clf,
                             feature_names=['Age', 'Gender', 'VitD_levels',
      ↪ 'Doc_visits', 'vitD_supp', 'Initial_admin', \
      ↪ 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight',
      ↪ 'Arthritis', \
      ↪ 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety',
      ↪ 'Allergic_rhinitis', \
      ↪ 'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',
      ↪ 'TotalCharge', \
      ↪ 'Additional_charges'], out_file=f)
```

```
[35]: predictions = clf.predict(test[['Age', 'Gender', 'VitD_levels', 'Doc_visits',
      ↪ 'vitD_supp', 'Initial_admin', \
      ↪ 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight',
      ↪ 'Arthritis', \
      ↪ 'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety',
      ↪ 'Allergic_rhinitis', \
      ↪ 'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days',
      ↪ 'TotalCharge', \
      ↪ 'Additional_charges']])
```

```
[36]: from sklearn.metrics import accuracy_score
accuracy_score(test['ReAdmis'], predictions)
```

```
[36]: 0.9755
```

0.1.10 Random Forest

```
[37]: n_estimators_selected = 2000
```

```
[38]: from sklearn.ensemble import RandomForestClassifier
predictors = ['Age', 'Gender', 'VitD_levels', 'Doc_visits', 'vitD_supp',
      ↪ 'Initial_admin', \
      ↪ 'HighBlood', 'Stroke', 'Complication_risk', 'Overweight',
      ↪ 'Arthritis', \
```

```

        'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety', \
        ↪ 'Allergic_rhinitis', \
        'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days', \
        ↪ 'TotalCharge', \
        'Additional_charges']
outcome = 'ReAdmis'

X = pruned_df[predictors]
y = pruned_df[outcome]

rf = RandomForestClassifier(n_estimators=n_estimators_selected, random_state=1, \
        ↪ oob_score=True)
rf.fit(X,y)
print(rf.oob_decision_function_)

```

```

[[0.9957204  0.0042796 ]
 [0.99463087 0.00536913]
 [0.99862448 0.00137552]
 ...
 [0.00551724 0.99448276]
 [0.00136054 0.99863946]
 [0.00416089 0.99583911]]

```

```

[39]: n_estimator = list(range(100, 3000, 100)) # Looks to flatten out @1500
oobScores = []
for n in n_estimator:
    rf = RandomForestClassifier(n_estimators=n,
                               criterion='entropy', max_depth=5,
                               random_state=1, oob_score=True)

    rf.fit(X, y)
    oobScores.append(rf.oob_score_)
    print(n) # quick and dirty sense of time

```


```

100
200
300
400
500
600
700
800
900
1000
1100
1200
1300
1400

```

1500
1600
1700
1800
1900
2000
2100
2200
2300
2400
2500
2600
2700
2800
2900

```
[40]: pd.DataFrame({  
      'n': n_estimator,  
      'oobScore': oobScores  
    }).plot(x='n', y='oobScore');
```



0.1.11 OOB Evaluation

```
[41]: # Import models and split utility function
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
# Set seed for reproducibility
SEED = 1
# Split data into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.3,
stratify= y,
random_state=SEED)
```

```
[42]: # Instantiate a classification-tree 'dt'
dt = DecisionTreeClassifier(max_depth=4,
min_samples_leaf=0.16,
random_state=SEED)
# Instantiate a BaggingClassifier 'bc'; set oob_score = True
bc = BaggingClassifier(base_estimator=dt, n_estimators=300,
oob_score=True, n_jobs=-1)
# Fit 'bc' to the training set
bc.fit(X_train, y_train)
# Predict the test set labels
y_pred = bc.predict(X_test)
```

```
/Users/jasonewilllis/opt/anaconda3/lib/python3.9/site-
packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was
renamed to `estimator` in version 1.2 and will be removed in 1.4.
warnings.warn(
```

```
[43]: # Evaluate test set accuracy
test_accuracy = accuracy_score(y_test, y_pred)
# Extract the OOB accuracy from 'bc'
oob_accuracy = bc.oob_score_
# Print test set accuracy
print('Test set accuracy: {:.3f}'.format(test_accuracy))
```

Test set accuracy: 0.977

```
[44]: # Print OOB accuracy
print('OOB accuracy: {:.3f}'.format(oob_accuracy))
```

OOB accuracy: 0.978

```
[45]: predictors = ['Age', 'Gender', 'VitD_levels', 'Doc_visits', 'vitD_supp',
↪ 'Initial_admin', \
```

```

        'HighBlood', 'Stroke', 'Complication_risk', 'Overweight', \
        ↪ 'Arthritis', \
        'Diabetes', 'Hyperlipidemia', 'BackPain', 'Anxiety', \
        ↪ 'Allergic_rhinitis', \
        'Reflux_esophagitis', 'Asthma', 'Services', 'Initial_days', \
        ↪ 'TotalCharge', \
        'Additional_charges']
outcome = 'ReAdmis'

X = pd.get_dummies(pruned_df[predictors], drop_first=True)
y = pruned_df[outcome]

rf_all = RandomForestClassifier(n_estimators=n_estimators_selected, \
    ↪ random_state=1)
rf_all.fit(X, y)

rf_all_entropy = RandomForestClassifier(n_estimators=n_estimators_selected, \
    ↪ random_state=1,
                                       criterion='gini')
print(rf_all_entropy.fit(X, y))

```

```
RandomForestClassifier(n_estimators=2000, random_state=1)
```

```

[46]: from collections import defaultdict
      from sklearn import metrics

      rf = RandomForestClassifier(n_estimators=n_estimators_selected)
      scores = defaultdict(list)

      # crossvalidate the scores on a number of different random splits of the data
      for _ in range(3):
          train_X, valid_X, train_y, valid_y = train_test_split(X, y,
                                                                  test_size=0.3)

          rf.fit(train_X, train_y)
          acc = metrics.accuracy_score(valid_y, rf.predict(valid_X))
          for column in X.columns:
              X_t = valid_X.copy()
              X_t[column] = np.random.permutation(X_t[column].values)
              shuff_acc = metrics.accuracy_score(valid_y, rf.predict(X_t))
              scores[column].append((acc-shuff_acc)/acc)
      print('Features sorted by their score:')
      print(sorted([(round(np.mean(score), 4), feat) for
                    feat, score in scores.items()], reverse=True))

```

Features sorted by their score:

```
[(0.4375, 'Initial_days'), (0.0088, 'TotalCharge'), (0.0008, 'Stroke'), (0.0005,
'Initial_admin'), (0.0005, 'Complication_risk'), (0.0005, 'Allergic_rhinitis'),
```

```
(0.0003, 'Asthma'), (0.0002, 'Services'), (0.0002, 'Overweight'), (0.0001,
'vitD_supp'), (0.0001, 'VitD_levels'), (0.0001, 'Gender'), (0.0, 'Anxiety'),
(0.0, 'Age'), (-0.0001, 'Hyperlipidemia'), (-0.0001, 'BackPain'), (-0.0001,
'Arthritis'), (-0.0001, 'Additional_charges'), (-0.0002, 'HighBlood'), (-0.0003,
'Diabetes'), (-0.0005, 'Reflux_esophagitis'), (-0.0006, 'Doc_visits')]
```

```
[47]: train_X, valid_X, train_y, valid_y = train_test_split(X, y, test_size=0.3)
      rf.fit(train_X, train_y)
      valid_X_predict = rf.predict(valid_X)
      acc = metrics.accuracy_score(valid_y, valid_X_predict)

      acc
```

```
[47]: 0.9793333333333333
```

```
[48]: set(rf.predict(valid_X))
```

```
[48]: {0, 1}
```

```
[49]: from sklearn.metrics import confusion_matrix

      conf_mat = confusion_matrix(valid_y, valid_X_predict)
      print(conf_mat)
      # TP, FP, FN, TN
```

```
[[1862   26]
 [   36 1076]]
```

```
[57]: #from sklearn.ensemble import RandomForestRegressor

      #regressor = RandomForestRegressor(n_estimators=20, random_state=0)
      #regressor.fit(train_X, train_y)
      #y_pred = regressor.predict(valid_X)

      #from sklearn.metrics import classification_report, confusion_matrix,
      ↪accuracy_score

      #print(confusion_matrix(y_test, y_pred))
      #print(classification_report(y_test, y_pred))
      #print(accuracy_score(y_test, y_pred))
```

```
[51]: from sklearn.metrics import classification_report, confusion_matrix,
      ↪accuracy_score

      print(confusion_matrix(y_test, y_pred))
      print(classification_report(y_test, y_pred))
      print(accuracy_score(y_test, y_pred))
```

```
[[1862  37]
 [ 33 1068]]
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	1899
1	0.97	0.97	0.97	1101
accuracy			0.98	3000
macro avg	0.97	0.98	0.97	3000
weighted avg	0.98	0.98	0.98	3000

```
0.9766666666666667
```

```
[52]: importances = rf_all.feature_importances_

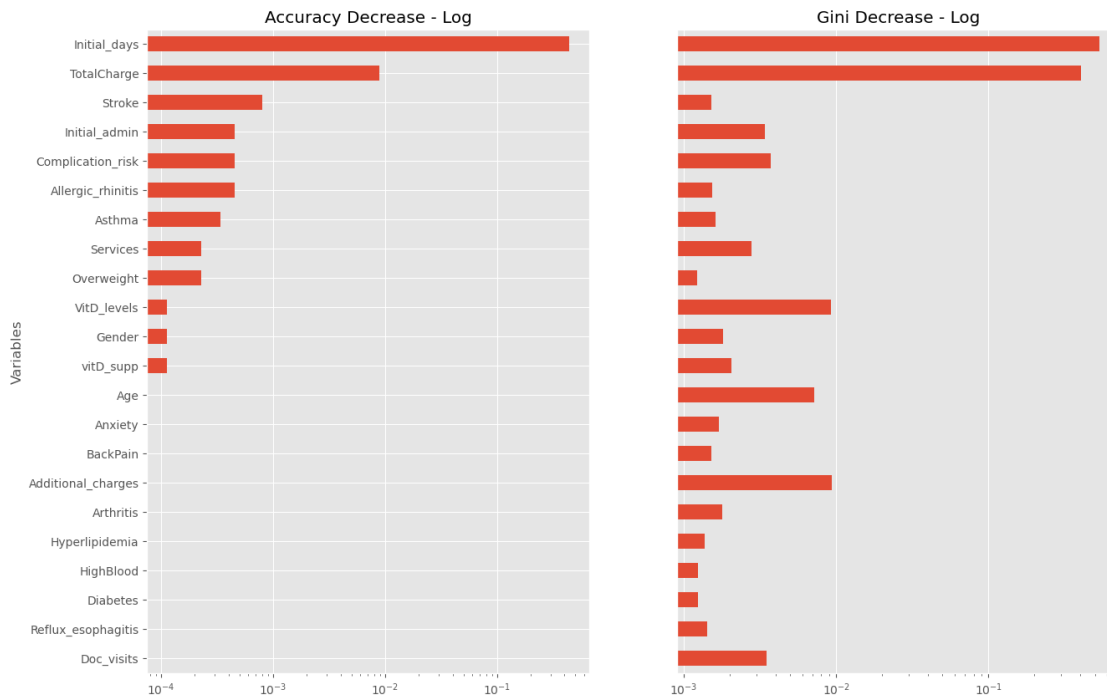
df = pd.DataFrame({
    'feature': X.columns,
    'Accuracy decrease': [np.mean(scores[column]) for column in
                          X.columns],
    'Gini decrease': rf_all.feature_importances_,
    'Entropy decrease': rf_all_entropy.feature_importances_,
})
df = df.sort_values('Accuracy decrease')

fig, axes = plt.subplots(ncols=2, figsize=(8, 5))
ax = df.plot(kind='barh', x='feature', y='Accuracy decrease', logx=True,
               figsize=(15,10), title='Accuracy Decrease - Log',
               legend=False, ax=axes[0])

ax.set_ylabel('Variables')

ax = df.plot(kind='barh', x='feature', y='Gini decrease', logx=True,
               title='Gini Decrease - Log',
               legend=False, ax=axes[1])
ax.set_ylabel('')
ax.get_yaxis().set_visible(False)

#plt.tight_layout()
plt.show()
```

0.1.12 Accuracy and MSE

```
[53]: from sklearn.metrics import mean_squared_error as MSE
      from sklearn.model_selection import cross_val_score
```

```
[55]: from sklearn import metrics
      print('Mean Absolute Error:', metrics.mean_absolute_error(train_y, valid_y))
      print('Mean Squared Error:', metrics.mean_squared_error(train_y, valid_y))
      print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(train_y,
      ↪valid_y)))
```

ValueError Traceback (most recent call last)

Cell In[55], line 2

```
1 from sklearn import metrics
----> 2 print('Mean Absolute Error:',
      ↪metrics.mean_absolute_error(train_y, valid_y))
3 print('Mean Squared Error:', metrics.mean_squared_error(train_y,
      ↪valid_y))
4 print('Root Mean Squared Error:', np.sqrt(metrics.
      ↪mean_squared_error(train_y, valid_y)))
```

File ~/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/_regression.py
↪196, in mean_absolute_error(y_true, y_pred, sample_weight, multioutput)

```

141 def mean_absolute_error(
142     y_true, y_pred, *, sample_weight=None, multioutput="uniform_average
143 ):
144     """Mean absolute error regression loss.
145
146     Read more in the :ref:`User Guide <mean_absolute_error>`.
147     (...)
148     0.85...
149     """
--> 196     y_type, y_true, y_pred, multioutput = _check_reg_targets(
197         y_true, y_pred, multioutput
198     )
199     check_consistent_length(y_true, y_pred, sample_weight)
200     output_errors = np.average(np.abs(y_pred - y_true),
↳ weights=sample_weight, axis=0)

File ~/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/_regression.py
↳ 100, in _check_reg_targets(y_true, y_pred, multioutput, dtype)
    66 def _check_reg_targets(y_true, y_pred, multioutput, dtype="numeric"):
    67     """Check that y_true and y_pred belong to the same regression task.
    68
    69     Parameters
    70     (...)
    71     correct keyword.
    72     """
--> 100     check_consistent_length(y_true, y_pred)
    101     y_true = check_array(y_true, ensure_2d=False, dtype=dtype)
    102     y_pred = check_array(y_pred, ensure_2d=False, dtype=dtype)

File ~/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py:
↳ 397, in check_consistent_length(*arrays)
    395 uniques = np.unique(lengths)
    396 if len(uniques) > 1:
--> 397     raise ValueError(
    398         "Found input variables with inconsistent numbers of samples: %r
    399         % [int(1) for 1 in lengths]
    400     )

ValueError: Found input variables with inconsistent numbers of samples: [7000,
↳ 3000]

```

0.1.13 Sharding and Exporting Data

```

[ ]: print("----- Prune Data -----")
print("pruned_df: ", pruned_df.shape)
print("")

```

```

print("----- Decision Tree Data -----")
print("train: ", train.shape)
print("test: ", test.shape)
print("")
print("----- Random Forest Data -----")
print("train_X: ", train_X.shape)
print("valid_X: ", valid_X.shape)
print("train_y: ", train_y.shape)
print("valid_y: ", valid_y.shape)

```

```

[ ]: # Export to CSV
pruned_df.to_csv('final_cleaned_dataset.csv', index=False)
#
train.to_csv('d_tree_train.csv', index=False)
test.to_csv('d_tree_test.csv', index=False)
#
train_X.to_csv('r_forest_train_X.csv', index=False)
valid_X.to_csv('r_forest_valid_X.csv', index=False)
train_y.to_csv('r_forest_train_y.csv', index=False)
valid_y.to_csv('r_forest_valid_y.csv', index=False)

```

```

[58]: # Import DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier
# Import train_test_split
from sklearn.model_selection import train_test_split
# Import accuracy_score
from sklearn.metrics import accuracy_score
# Split the dataset into 80% train, 20% test
X_train, X_test, y_train, y_test= train_test_split(X, y,
test_size=0.2,
stratify=y,
random_state=1)
# Instantiate dt
dt = DecisionTreeClassifier(max_depth=2, random_state=1)

```

```

[59]: # Fit dt to the training set
dt.fit(X_train,y_train)
# Predict the test set labels
y_pred = dt.predict(X_test)
# Evaluate the test-set accuracy
accuracy_score(y_test, y_pred)

```

[59]: 0.9745

```

[60]: # Import DecisionTreeClassifier
from sklearn.tree import DecisionTreeClassifier
# Import train_test_split

```

```

from sklearn.model_selection import train_test_split
# Import accuracy_score
from sklearn.metrics import accuracy_score
# Split dataset into 80% train, 20% test
X_train, X_test, y_train, y_test= train_test_split(X, y,
test_size=0.2,
stratify=y,
random_state=1)
# Instantiate dt, set 'criterion' to 'gini'
dt = DecisionTreeClassifier(criterion='gini', random_state=1)

```

```

[61]: # Fit dt to the training set
dt.fit(X_train,y_train)
# Predict test-set labels
y_pred= dt.predict(X_test)
# Evaluate test-set accuracy
accuracy_score(y_test, y_pred)

```

[61]: 0.967

```

[62]: # Import DecisionTreeRegressor
from sklearn.tree import DecisionTreeRegressor
# Import train_test_split
from sklearn.model_selection import train_test_split
# Import mean_squared_error as MSE
from sklearn.metrics import mean_squared_error as MSE
# Split data into 80% train and 20% test
X_train, X_test, y_train, y_test= train_test_split(X, y,
test_size=0.2,
random_state=3)
# Instantiate a DecisionTreeRegressor 'dt'
dt = DecisionTreeRegressor(max_depth=4,
min_samples_leaf=0.1,
random_state=3)

```

```

[63]: # Fit 'dt' to the training-set
dt.fit(X_train, y_train)
# Predict test-set labels
y_pred = dt.predict(X_test)
# Compute test-set MSE
mse_dt = MSE(y_test, y_pred)
# Compute test-set RMSE
rmse_dt = mse_dt**(1/2)
# Print rmse_dt
print(rmse_dt)

```

0.14757422655650276

```
[64]: from sklearn.ensemble import RandomForestRegressor

SEED = 1
X_train, X_test, y_train, y_test= train_test_split(X, y,
test_size=0.3,
random_state=SEED)

rf = RandomForestRegressor(n_estimators=400,
                           min_samples_leaf=0.12,
                           random_state=SEED)

rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)
```

```
[65]: rmse_test = MSE(y_test, y_pred)**(1/2)

print('Test set RMSE of rf: {:.2f}'.format(rmse_test))
```

Test set RMSE of rf: 0.14

```
[66]: importances_rf = pd.Series(rf.feature_importances_, index = X.columns)

# Sort
sorted_importances_rf = importances_rf.sort_values()

# Horizontal Bar Plot
sorted_importances_rf.plot(kind='barh', color='lightgreen'); plt.show()
```

